

Assignment 3

CS4533 Fall 2020

Due: Sunday 11/29 at 11:55pm. Total Score: 165 points

Note: This assignment has 4 pages.

This assignment extends Assignment 2 by adding the option of producing the *shadow effect* (with respect to a fixed light source) while rolling a sphere, as well as the options of shading and lighting. You will practice implementing three main techniques: (1) producing a shadow, (2) drawing a *decal* on top of an object — in this case the shadow is a decal on top of the ground, and (3) shading and lighting using programmable shaders. (Consult Lecture Notes for techniques of producing shadow and decal effects, and Lecture Notes and sample code “Handout: rotate-cube-shading.cpp” for shading and lighting.)

General Instructions: same as Assignment 1.

(a) Consider the two unit-sphere files given in Assignment 2 (<http://cse.poly.edu/cs653/assg2/>). There are two more unit-sphere files given in <http://cse.poly.edu/cs653/assg3/>, with 256 and 1024 triangles, respectively. The file format is described in Assignment 2.

Draw the x-, y-, z-axes, the “ground” (the quadrilateral indicating the x-z plane), and roll the wire-frame sphere as in Assignment 2. Your task here is to add the corresponding shadow of the sphere on the ground given a fixed light source located at position $L = (-14.0, 12.0, -3.0, 1.0)$. Draw the shadow with color $(0.25, 0.25, 0.25, 0.65)$. Note that the shadow should be the wire-frame sphere projected onto the ground, projected from the light source. As the sphere rolls, the shadow should also “roll” accordingly.

As another option, draw the sphere as a *solid* sphere and roll it as before, by drawing the sphere triangles as *filled* triangles (by setting a suitable mode in `glPolygonMode()`). Produce the corresponding shadow effect for this rolling solid sphere as well, using the same shadow color and the same light source location L . (The options are put together by a menu; see **Part (b)**.)

Hint: Draw the shadow as a separate object. As discussed in class, the shadow object is obtained as follows. At each frame, from each generic point $p = (x, y, z, 1)$ on the sphere (the sphere is already **at its final position and orientation at this frame**), project p onto the ground using the ray going from L through p and intersecting the ground plane at point q ; q is then the shadow point of p . Since p can be any point on the sphere, any corresponding shadow point q can be obtained this way. Suppose $p = (x, y, z, 1)$ and $q = (x_q, y_q, z_q, 1)$. Derive x_q and z_q in terms of x, y , and z , using simple geometry and the fact that the ground plane is $y = 0$ and thus $y_q = 0$. Then, derive a 4×4 transformation matrix N to carry out this “shadow projection” that transforms p to q using the homogeneous-coordinate representation.

Note that in OpenGL, the homogeneous-coordinate representation (x, y, z, w) , $w \neq 0$, is equivalent to the representation $(x/w, y/w, z/w, 1)$, i.e., the *perspective division* to convert (x, y, z, w) to $(x/w, y/w, z/w, 1)$ is *automatically done* by OpenGL. Also, to produce the shadow, you first perform the sphere-rolling transformation (as in Assignment 2) on the input sphere, and then perform the shadow projection transformation N .

Programming Tips:

In OpenGL, although the perspective division is automatically done to convert (x, y, z, w) to $(x/w, y/w, z/w, 1)$ for $w \neq 0$, the OpenGL implementation actually requires that $w > 0$. So if you tried $q = (\frac{f(x,y,z)}{h(x,y,z)}, 0, \frac{g(x,y,z)}{h(x,y,z)}, 1) \equiv (f(x,y,z), 0, g(x,y,z), h(x,y,z))$ and it did not work, try to use $q = (\frac{-f(x,y,z)}{-h(x,y,z)}, 0, \frac{-g(x,y,z)}{-h(x,y,z)}, 1) \equiv (-f(x,y,z), 0, -g(x,y,z), -h(x,y,z))$. **(40 points)**

(b) If you just draw the shadow as in **Part (a)**, then you will see a *broken* shadow. The undesirable effect is especially obvious for the solid sphere. This is because the shadow and the ground are on the same plane, but the z-buffer only has a limited numerical precision. As a result, some portion of the shadow may appear in front of the ground and some portion may appear *behind* the ground and thus *blocked* by the ground; the result is unpredictable. What you want here is to make the shadow *always appear on top of* the ground, i.e., to make the shadow a *decal* on top of the ground. Use the technique of making a decal, as discussed in class, to achieve the desired effect.

Recall from Assignment 2 that there is a menu associated with the left mouse button. Add a menu entry “Shadow”, and implement it as a submenu with 2 submenu entries:

“No” — do not produce the shadow, as in Assignment 2, and

“Yes” — produce the *correct* shadow, as implemented in **Parts (a) and (b)**.

For the same menu, add menu entries “Enable Lighting”, “Wire Frame Sphere”, and “Shading”. Implement “Enable Lighting” as a submenu with 2 submenu entries:

“No” — the lighting is disabled, as is the case so far, and

“Yes” — the lighting is enabled, to be implemented in **Parts (c) and (d)**.

Also, when “Wire Frame Sphere” is selected, the sphere should be drawn as a wire-frame sphere (as in Assignment 2).

In addition, implement “Shading” as a submenu with 2 submenu entries: “flat shading” and “smooth shading”. When lighting is disabled (by choosing “Enable Lighting” and then “No”), *both* options should just draw the sphere as a solid sphere (see **Part (a)**) and give the same result. When lighting is enabled, “flat shading” and “smooth shading” will be further implemented in **Part (c)**.

Note that choosing “Wire Frame Sphere” should draw a wire-frame sphere, and choosing both “flat shading” and “smooth shading” should draw a solid sphere; in all these cases you should also produce the corresponding *correct* shadow if “Shadow” is “Yes”. The ground should be rendered according to whether “Enable Lighting” is “Yes” or “No”, *independent* of “Wire Frame Sphere”. Namely, you always render the solid ground. When “Enable Lighting” is “No”, the ground is rendered with a single given color as in Assignment 2. When “Enable Lighting” is “Yes”, the ground is rendered with the shading computation performed.

Hint:

1. You should always enable the z-buffer testing (by `glEnable(GL_DEPTH_TEST)`). Also, you need to draw the ground/shadow in multiple passes, enabling/disabling writing to the z-buffer or color buffer (color buffer is the frame buffer), as discussed in class.
2. Use `glDepthMask(GL_FALSE)` to disable writing to the z-buffer, `glDepthMask(GL_TRUE)` to enable writing to the z-buffer, `glColorMask(GL_FALSE, GL_FALSE, GL_FALSE, GL_FALSE)` to disable writing to the color buffer (in RGBA color mode), and `glColorMask(GL_TRUE, GL_TRUE, GL_TRUE, GL_TRUE)` to enable writing to the color buffer. **(40 points)**

In the remaining parts (**Parts (c) and (d)**), your task is to add shading and lighting effects when “Enable Lighting” is “Yes”.

Note: You should use **programmable shaders** to implement shading and lighting. The OpenGL commands for fixed-functionality shading and lighting are deprecated and **not allowed**.

(c) Provide a **global ambient light** with white color (1.0, 1.0, 1.0, 1.0), and also provide a **directional (distant) light source** with black ambient color (0.0, 0.0, 0.0, 1.0), diffuse color (0.8, 0.8, 0.8, 1.0), specular color (0.2, 0.2, 0.2, 1.0), and direction (0.1, 0.0, -1.0, 0.0) in the **eye coordinate system** (i.e., the **right-handed** coordinate system where the eye is at the origin looking toward the $-z$ direction — this is the coordinate system after applying the `LookAt()` transformation to the world coordinate system). Note that (0.1, 0.0, -1.0, 0.0) is a vector and this light source is a distant (directional) light.

In addition, give your background and the x-, y-, and z-axes the same colors as in Assignment 2. Give your quadrilateral “ground” that indicates the x-z plane a green diffuse color (0.0, 1.0, 0.0, 1.0), with ambient color (0.2, 0.2, 0.2, 1.0) and specular color (0.0, 0.0, 0.0, 1.0), and give your sphere a golden yellow diffuse and specular color (1.0, 0.84, 0.0, 1.0), with ambient color (0.2, 0.2, 0.2, 1.0) and a shininess coefficient of 125.0.

Consider the four unit-sphere files given. Actually, there is one more piece of information implicitly given in the file format: each triangle (given by the coordinates of its three vertices) has its three vertices ordered such that if the fingers of the right hand are curled along the direction of the vertex specification, the thumb will point towards the triangle’s *outward normal* (in this case this normal points towards the *outside* of the sphere). Augment your data structure from Assignment 2 so that associated with each triangle you also store its outward normal vector of *length 1* (i.e., the *unit-length* outward normal vector) in *each of the 3 triangle vertices*.

Recall that the menu entry “Shading” has two submenu entries “flat shading” and “smooth shading”. Implement these entries with the following shading effects when “Enable Lighting” is “Yes”:

- (1) “flat shading” — flat shade the *filled* triangles of the unit sphere, where the unit-length normal vector associated with each triangle is the normal vector that you just computed and stored.
- (2) “smooth shading” — smooth shade the *filled* triangles of the unit sphere. To assign the unit-length vertex normals, use the fact that if the unit sphere is centered at the origin o and v is a point on the surface of the unit sphere, then \vec{ov} is the unit-length normal vector at v .

Notice that the options for the wire-frame sphere and the shadow effects should still work, as you worked out before. (For the wire-frame sphere, do not perform lighting/shading computation but instead directly use the color as the case where “Enable Lighting” is “No”.)

Hint: To perform lighting/shading computation, everything you draw needs normal vector(s). In particular, you need to specify the **ground normal for the ground**, and the **normals for the sphere as described in (1) and (2)** above.

However, when you draw the 3 axes, the wire-frame sphere, and/or the sphere shadow, you should not perform lighting/shading computation. Instead, directly use the colors of them as the case where “Enable Lighting” is “No”. **(45 points)**

(d) Provide **another light source** which is **positional**, with white diffuse and specular color (1.0, 1.0, 1.0, 1.0), black ambient color (0.0, 0.0, 0.0, 1.0), with position at $L = (-14.0, 12.0, -3.0, 1.0)$

(in the usual **world coordinate system**). This is the light source that produces the shadow. Set up its constant, linear, and quadratic attenuation values to be 2.0, 0.01, and 0.001, respectively. Give two choices for the type of this light source:

- (1) spot light, whose direction is from its position toward the point $(-6.0, 0.0, -4.5, 1.0)$ (again in the usual **world coordinate system**), with the exponent value 15.0 and the cutoff angle 20.0° , and
- (2) point source.

In the same menu (associated with the left mouse button), add a new menu entry called “Light Source”. Implement it as a submenu with two submenu entries: “spot light” (corresponding to (1) above) and “point source” (corresponding to (2) above). Your program should allow the user to switch between (1) and (2). **(40 points)**

Useful Tips:

1. For the function `cos (angle)` in GLSL, the parameter `angle` must be in **radians** (rather than in degrees).
2. Typically shading/lighting computation is done in the **eye coordinate system** since it is more convenient to do it this way. For light position and direction, since they do not change with respect to different vertices (or fragments), they are first transformed to the **eye coordinate system** and then passed on to the shaders as **uniform variables**.
3. Although there are multiple shading options, multiple light sources (a directional light source in **Part (c)** and a positional light source in **Part (d)**), and multiple objects (sphere, ground, 3 axes), the basic shading computations for them are the same or similar. It is advised that you use only **one pair** of shaders (one vertex shader and one fragment shader) for all these tasks. Write a common function in the (vertex) shader that can be used by different light sources for the basic shading computation, and use flags (passed on to the shaders as **uniform variables**) to control which part of the shader code to execute for the different shading options, light sources, and/or objects.